ABHISHEK SUKHADIYA

San Francisco Bay Area (Open to relocate)

+1 510-980-1735 | Linkedin\Abhishek-Sukhadiya | Github\Abhishek-Sukhadiya | abhisheksukhadiyaus38@gmail.com

SUMMARY

Software Engineer (MS CS) specializing in **Agentic AI** and **LangGraph** workflows. Expert in building scalable **LLM-based** solutions, **prompt engineering**, and high-concurrency **Python** backends with a focus on **stateful graph design**.

TECHNICAL SKILLS

AI/ML: LangGraph, Hugging Face, LLMs (GPT-4, Claude, Llama), Prompt Engineering, RAG, Agent Orchestration, PyTorch, Azure OpenAI
Languages: Python (Advanced/Async), Java, Javascript, C++, SQL, Shell Scripting, HTML, CSS, PHP, NoSQL
Backend & Cloud: REST APIs, FastAPI, Docker, CI/CD (Jenkins), AWS (S3), Linux/Unix, Azure AI Foundry, Azure Functions, Datadog
Data Engineering: PostgreSQL, MySQL, MongoDB, Vector Databases (Pinecone/Chroma), Schema Ingestion & performance Tuning

PROFESSIONAL EXPERIENCE

Research Associate Scientist at San Francisco Bay University | Aug 2025 – Dec 2025

- Developed RAG-based LLM solutions, optimizing prompt techniques to achieve high accuracy despite strict GPU memory limitations
- Mentored 50+ students on SQL optimization, ensuring robust database designs under the constraint of a 48-hour grading turnaround

AI & Software Engineer Intern | BeProEx & Vtechys | Oct 2024 – Aug 2025

- Architected LangGraph agent orchestrations for new features, optimizing stateful workflows to meet sub-second latency requirements
- Refined prompt engineering for high-fidelity response alignment while adhering to strict safety guardrails and token cost limits
- Designed evaluation harnesses for iterative GenAI testing, enabling rapid validation despite limited availability of ground truth
- Built modular Python automation and SQL/NoSQL services, cutting maintenance by 30% while integrating with legacy system constraints

Software & AI Engineer, Backend | Musikaar (Vendor for Tenable) | Jan 2021 – Dec 2023

- Developed scalable Python backends for data ingestion, managing 100M+ records within strict MySQL storage and memory constraints
- Optimized SQL queries to slash latency by 50%, maintaining sub-second response times despite high-volume concurrent traffic
- Architected Python/Java ingestion for diverse partner schemas, boosting throughput 70% to meet strict 24-hour data delivery SLAs
- Resolved 30+ critical C++ SDK issues, ensuring 99.9% reliability in a distributed system under high-availability requirements
- Automated validation with Python, cutting testing time by 60% to accelerate release cycles despite limited manual QA resources
- Fine-tuned LLMs for automated threat classification, achieving 92% accuracy while operating within strict 16GB GPU memory limits

Machine learning & Backend Engineer| NetAnalytics | Jul 2020 – Nov 2020

- Automated network discovery via Python/SSH, managing 1k+ nodes despite inconsistent legacy firmware and rigid firewall constraints
- Built secure REST APIs and Prometheus dashboards for real-time monitoring under strict data privacy and minimal infrastructure budget

SOFTWARE & RESEARCH PROJECTS

Gemini CLI Turbo (Enterprise-Grade LLM Automation & Performance Optimization)

- Built an enterprise-grade Python CLI for LLM APIs, cutting latency 60% while managing strict rate limits and high token costs
- Implemented fault-tolerance via circuit breakers and semantic caching to ensure 99.9% uptime despite unstable third-party APIs
- Designed a scalable Docker-based architecture for large AI workloads, optimizing resource utilization to cut cloud infra costs

Digital Sage (Gujarati RAG System)

- Developed a Gujarati RAG system using LangChain, overcoming severe data scarcity in low-resource languages via custom embeddings
- Achieved 95% accuracy using citation-first prompts and eval harnesses, eliminating hallucinations for sensitive enterprise data
- Built an OCR and normalization pipeline for semantic search, processing non-standard document formats under strict time constraints

Multi-Agent Orchestration Framework (Advanced LangGraph)

- Engineered a LangGraph multi-agent system with Checkpointers, maintaining persistent state under tight memory and compute limits
- Orchestrated Coder/Reviewer agents, boosting accuracy 30% via iterative loops while optimizing for strict LLM token budget limits
- Designed stateful graph workflows for complex paths, ensuring reliable enterprise automation for 1,000+ concurrent user sessions

EDUCATION & CERTIFICATIONS

MS in Computer Science | San Francisco Bay University, CA | Graduated Dec 2025

- GPA: 3.82 | Coursework: AI and ML, Advanced Algorithms, Software Architecture, Databases

BS in Information Technology | Silver Oak College of Engineering & Technology | Graduated 2020

Microsoft Certified: Azure AI Engineer Associate (AI-102) | 2025